

XAVIER GOAS AGUILILLA

A CORPUS-BASED APPROACH TO
THE LANGUAGE OF THE PAPYRI



Master's thesis submitted in partial fulfilment of
the requirements for the degree of

MASTER OF ARTS IN DE TAAL-
EN LETTERKUNDE

Supervisor: dr. T. Van Hal
KU Leuven
Faculty of Arts
Department of Greek Studies

LEUVEN, 2012

Xavier Goas Aguililla: *A corpus-based approach to the language of the papyri*, © June 2012.

E-MAIL:

xavier.goas@student.kuleuven.be

ABSTRACT

SAMENVATTING

Deze masterpaper bevat **25383** tekens.

PREFACE

The papyri are an invaluable resource for documenting the history and evolution of the Greek language. The recently published *The Language of the Papyri* [Evans and Obbink, 2010] has placed the spotlight firmly on the potential of this field while at the same time pointing out the regrettable lack of recent scholarly work available in it.

Though there is now a collection of texts available that is well-formatted and can easily be converted into an adequate corpus for linguistic research, comparatively few scholars are interested in exploiting the available resources; and none, to my knowledge, have attempted to do so.

A first necessity is, of course, a broad *status quaestionis* yet in a less traditional sense than one might expect: the focus here lies more upon technical and technological concerns than specialised monographs for which data was gathered manually. Nevertheless, I have for the sake of completeness chosen to include a comprehensive bibliography of the field for clear reference. The term bibliography would be rather less appropriate in describing the array of databases and linguistic tools available to us - it is rather a select catalogue documenting the most relevant items in the instrumentarium in some detail while providing a briefly annotated list of other useful resources.

A second chapter is dedicated

1. a survey of previous studies on the language of the Greek papyri and an analysis of the methodology used therein;
2. a critical evaluation of the methods used by aforementioned studies, with equal attention for both positive and negative aspects of the applied method;
3. the development of a methodology based on this evaluation that is fit for use in corpus-based grammatical and linguistic studies;
4. an analysis and evaluation of the available tools in the field of papyrology for such studies, followed by suggestions for possible improvements to these tools;
5. the study of select grammatical and linguistic questions concerning the language of the Greek papyri using the methodology and toolset developed previously — essentially a practical test of the previous findings.

ACKNOWLEDGEMENTS

CONTENTS

PREFACE [v](#)

ACKNOWLEDGEMENTS [v](#)

1	INTRODUCTION AND PRELIMINARIES	1
1.1	Thesis	1
1.2	Preliminaries	1
1.2.1	The language of the papyri	1
1.2.2	Corpus linguistics	3
1.2.3	The digital classics	4
1.2.4	Natural language processing	6
1.3	Methodology	7
2	THE CORPUS	9
2.1	Goals	9
2.2	Design	10
2.2.1	Initial inquiries	10
2.3	Technical	12
2.3.1	Requirements	12
2.3.2	Extracting training data	12
2.3.3	Principles of statistical natural language processing	12
2.3.4	Accuracy testing	12
3	THE TOOL	13
3.1	Goals	13
3.2	Design	13
3.3	Technical	13
3.4	Interface	14
4	APPLICATIONS	15
4.1	Collaborative editing	15
4.2	Corpus-based grammars	15
4.3	Textual criticism	15
5	CONCLUSION	17
	BIBLIOGRAPHY	19

1

INTRODUCTION AND PRELIMINARIES

The study of the language of the papyri has in the past thirty years seen little evolution until the recent appearance of Evans and Obbink's *The Language of the Papyri* [Evans and Obbink, 2010], which has placed the subject in the spotlight again. Twentieth-century scholarship on the topic, though still useful for those interested in the study of the papyri for historical purposes, is either antiquated, limited in scope or incomplete (see *infra*). Despite this, the papyri are useful source material for the history and evolution of the Greek language, as they contain not only official texts but private documents as well, whose linguistic features and peculiarities have the potential to foster new insights into the nature of colloquial Greek.

1.1 THESIS

The following thesis intends to prove that it is possible to generate basic linguistic annotation for a large digitalised corpus of papyri in ancient Greek using readily available tools and techniques with minimal technical overhead. Such a corpus could be a boon to scholars interested in the Greek of the papyri, as it would facilitate, for instance, the creation of linguistically sound grammars and lexica.

1.2 PRELIMINARIES

The following section will provide a background sketch, consisting of a short overview of previous efforts and an elucidation of some key concepts.

1.2.1 The language of the papyri

The papyri began to be studied linguistically not by papyrologists and historians, but rather by Bible scholars and grammarians interested in their relevance in the development of koinê Greek, particularly that of the New Testament. G. N. Hatzidakis, W. Crönert, K. Dieterich, A. Deissmann, and A. Thumb pioneered the field in the late nineteenth and early twentieth century, spurring a resurgence of scholarship on the topic; an excellent overview of pre-1970s research may be found in Mandilaras [1973] and Gignac [1976 and 1981].

During this period, Mayser began work on the earliest compendious grammar of the papyri; it limits itself to the Ptolemaic era but

explores it at length and in great detail. The work consists of a part on phonology and morphology, made up of three slimmer volumes, and a part on syntax, encompassing three larger volumes. Its composition seems to have been exhausting: it took Mayser thirty-six years to finish volumes I.2 through II.3, with I.1 only completed in 1970 by Hans Schmoll, at which point the entire series was given a second edition.

When casually browsing through some of its chapters (though casual is hardly the word one would associate with the *Grammatik*) it is remarkable to see that Mayser brings an abundance of material to the table for each grammatical observation he makes, however small it may be. For instance, the section on diminutives essentially consists of pages upon pages of examples categorised by their endings.

This is its great strength as a reference work - whenever one is faced with an unusual grammatical phenomenon in any papyrus, consulting Mayser is bound to clarify the matter; or rather, it was, for the work is now inevitably dated. The volumes published during Mayser's lifetime only include papyri up to their date of publication; only the first tome by Schmoll includes papyri up to 1968. It is still a largely useful resource, but it is in urgent need of refreshment.

After Mayser set the standard for the Ptolemaic papyri, a grammar of the post-Ptolemaic papyri was the new *desideratum* in papyrology. The work had been embarked on by Salonijs, Ljungvik, Kapsomenos, and Palmer, only to be interrupted or thwarted by circumstance or lack of resources. Salonijs [1927], for instance, only managed to write an introduction on the sources, though he offered valuable comments on the matter of deciding how close to spoken language a piece of writing is. Ljungvik [1932] contains select studies on some points of syntax.

It is in the 1930's that we see attempts to create a grammar of the papyri that would be the equivalent of Mayser for the post-Ptolemaic period. Kapsomenos published a series of critical notes [1938, 1957] on the subject; though he attempted at a work on the scale of the *Grammatik*, he found the resources sorely lacking, as the existing editions of papyrus texts could not form the basis for a systematic grammatical study. The other was Palmer, who had embarked on similar project and had already set out a methodology [1934]; the war interrupted his efforts, and he published what he had already completed, a treatise on the suffixes in word formation [1945].

A new work of some magnitude presents itself two decades later with B. G. Mandilaras' *The verb in the Greek non-literary papyri* [1973]. Though it does not aim to be a grammar of the papyri, it does offer a thorough and satisfactory treatment of the verbal system as manifest in the papyri. Further efforts essentially do not appear until the publication of Gignac's grammar. It is essentially treading in the footsteps of Mayser, only with further methodological refinement and a more limited, though still sufficiently exhaustive, array of examples. The author, for reasons unknown to me, only managed to complete two of the three projected volumes, on phonology and on morphology. The volume on syntax is thus absent, a gap only partly filled by Mandilaras' *The verb in the Greek non-literary papyri*.

Finally, there is the aforementioned *The Language of the Papyri* [Evans and Obbink, 2010], which does not aim to be a work on the same scale as the aforementioned. It is a collection of articles on various topics, the whole of which is meant to illuminate new avenues for future research. A particularly relevant chapter for this thesis is the last one by Porter and O'Donnell [Porter and O'Donnell, 2010], who set out to create a linguistic corpus for a selection of papyri; their tagging approach, however, is manual, and their target corpus limited. The authors also are the creators of <http://www.opentext.org/>, a project aiming for the development of annotated Greek corpora and tools to analyse them; sadly, no progress seems to have been made since 2005.

1.2.2 Corpus linguistics

A¹ corpus or text corpus is a large, structured collection of texts designed for the statistical testing of linguistic hypotheses. The core methodological concepts of this mode of analysis may be found in the concordance, a tool first created by biblical scholars in the Middle Ages as an aid in exegesis. Among literary scholars, the concordance also enjoyed use, although to a lesser degree; the eighteenth century saw the creation of a concordance to Shakespeare.

The development of the concordance into the modern corpus was not primarily driven by the methods of biblical and literary scholars; rather, lexicography and pre-Chomskyan structural linguistics played a crucial role.

Samuel Johnson created his famous comprehensive dictionary of English by means of a manually composed corpus consisting of countless slips of paper detailing contemporary usage. A similar method was used in the 1880s for the Oxford English Dictionary project - a staggering three million slips formed the basis from which the dictionary was compiled.

1950s American structuralist linguistics was the other prong of progress; its heralding of linguistic data as a central given in the study of language supported by the ancient method of searching and indexing ensures its proponents may be called the forerunners of corpus linguistics.

Computer-generated concordances make their appearance in the late 1950s, initially relying on the clunky tools of the day - punch cards. A notable example is the *Index Thomisticus*, a concordance to the works of Thomas of Aquino created by the late Roberto Busa S.J. which only saw completion after thirty years of hard work; the printed version spans 56 volumes and is a testament to the diligence and industry of its author. The 1970s brought strides forward in technology, with the creation of computerised systems to replace catalogue indexing cards, a change that greatly benefited bibliography and archivistics.

It is only in the 1980s and 1990s that are marked the arrival of fully developed corpora in the modern sense of the word; for though the basic concepts of corpus linguistics were already widely used, they could

¹ The following section is based *passim* on McCarthy and O'Keeffe [2010].

not be applied on a large scale without the adequate tools. The rise of the desktop computer and the Internet as well as the seemingly ever-rising pace of technological development ensured the accessibility of digital tools. The old tools - punch cards, mainframes, tape recorders and the like - were gladly cast aside in favour of the new data carriers.

The perpetual increase of computing power equally demonstrated the limits of large-scale corpora; while lexicographical projects that had as their purpose to document the greatest number of possible usages could keep increasing the size of their corpora, the size of others went down as they whittled the data down to a specific set of uses of language.

The possible applications of the techniques of corpus linguistics are diverse and numerous; for they allow for a radical enlargement in scope while remaining empirical, and remove arduous manual labour from the equation. Corpus linguistics can be an end to itself; it can, however, assert an important role in broader research. McCarthy and O’Keeffe [2010, p. 7] mention areas such language teaching and learning, discourse analysis, literary stylistics, forensic linguistics, pragmatics, speech technology, sociolinguistics and health communication, among others.

The term ‘corpus’ has a slightly different usage in classical philology: they designate a structured collection of texts, but that collection is not primarily intended for the testing of linguistic hypotheses. Instead, we have, for instance, the ancient corpus Tibullianum, or modern-day collection, for instance the Corpus Papyrorum Judaicarum, etc. We are primarily interested in the digital techniques used to create linguistic corpora; so let us first take a look at the progress of the digital classics.

1.2.3 The digital classics

Classical philology, despite its status as one of the oldest and most conservative scientific disciplines still in existence today, has in the last fifty years found itself at the front lines of the digital humanities movement. Incipient efforts in the fifties and sixties, mainly stylometric and lexical studies and the development of concordances, demonstrated the relevance of informatics in the classics, an evolution that was at first met with some skepticism, but later fully embraced.

The efforts began with the aforementioned Index Thomisticus, the first computer-based corpus in a classical language; but the first true impetus was the foundation of the Thesaurus Linguae Graecae project in 1972, a monumental project with as its goal the stocking of all Greek texts from the Homeric epics to the fall of Constantinople. Over the years, many functions have been added to this ever more powerful tool; and even in the beginning stages of its development, the TLG garnered praise.

The usefulness of the tool in its current form cannot be overstated: not only does it contain a well-formatted and easily accessible gigantic collection of text editions whose scope and dimensions exceed those of nearly any university library; it also offers all of these texts in a format

that allows for lexical, morphological and proximity searches, as well as including a full version of the Liddell & Scott and Lewis & Short dictionaries. The TLG has become a staple of the digital classics.

Despite this, the TLG is becoming more and more dated as technology progresses. While recent years have seen the rise of Unicode as the standard for encoding ancient Greek, the TLG still uses beta code, a transliteration system designed to only use the ASCII character set, and the texts are stored using an obsolete text-streaming format from 1974, which divides the text in blocks of eight kilobytes and marks the division between segment.

A digitised version of the Liddell-Scott-Jones lexicon has been added to the TLG's web interface, but the texts themselves have not undergone extensive tagging, only lemmatisation. Searching through the database can be done by searching for specific forms of a lemma, or by searching for all forms of a lemma, but this is essentially the limit of the search tool's power; it is not possible to perform a query for all possible lemmata associated with a particular form, i.e. we cannot find all forms which are, for example, an active perfect indicative.

In the wake of the TLG, several notable projects have emerged: Brepols' Library of Latin Texts is trying hard to be for Latin texts what the TLG is for Greek texts; the Packard Humanities Institute has released CD's containing a selection of classical Latin works. In more recent times, the Perseus Project has enjoyed great popularity because of the attractive combination of an excellent selection of classical texts with translations, good accessibility and a set of interesting textual tools, the entire package carrying a very interesting price tag for the average user — it is free to use, and for the greatest part, open source as well.

The databases I have mentioned are quite general in scope; but within the domain of classical philology, other specialised projects exist. Within the field of papyrology, for instance, the digital revolution has taken a firm foothold. Starting with several separate databases, the field has experienced a tendency towards convergence and integration of the available resources, as exemplarised by the *papyri.info* website, maintained by Columbia University, that integrates the main papyrological databases into a single database.

A great feature of this database is the shell in which all data is wrapped; they are compliant with the EpiDoc standard, a subset of XML based on the TEI standard and developed specifically for epigraphical and papyrological texts. One may access the database's resources through the Papyrological Navigator and suggest corrections and readings through the Papyrological Editor. What's more, all data is freely accessible under the Creative Commons License, crowd-sourced, regularly updated, and can be downloaded for easier searching and tweaking.

In other words, *papyri.info* has brought the open-source mentality from the computer world into the classics. For our purposes, this open setup is desirable, as the database is not fit for them as it is, but can with some effort be molded into a useful tool.

1.2.4 Natural language processing

Natural language processing (henceforth NLP) is a subdiscipline in computer science concerned with the interaction between natural human language and computers. Its history well and truly starts in the fifties, with a basic concept which has played a great role in natural language processing, and computer science in general, the Turing test. This test, put forth by Alan Turing in his seminal paper *Computing Machinery and Intelligence* [Turing, 1950], evaluates whether a machine is intelligent or not by placing a human in conversation with another human and a machine; if the first human cannot tell the other human and the machine apart, the machine passes the test.

Machine translation systems entered development, though progress soon stalled because of technical limitations and because of methodological obstacles: such systems were dependent on complex rulesets written by programmers that allowed for very little flexibility. Because of the slow return on investments made, funding for artificial intelligence in general and machine translation specifically was drastically reduced throughout the late sixties and the seventies.

A resurgence followed: thanks to advances in computational power and the decline of Chomskyan linguistics, which had been the dominant theoretical vantage point in the preceding thirty years, the eighties were marked by the introduction of statistical machine translation, which is fundamentally based on the tenets of corpus linguistics. Modern natural language processing is therefore situated on the crossroads between various fields: artificial intelligence, computer science, statistics, and corpus and computational linguistics. It looks to be an exciting field for the coming years as its techniques are under constant improvement and ever more present in our daily lives.

Most NLP software is designed explicitly with living languages in mind; English, being a world language and the international *lingua franca*, has enjoyed most of the attention, but other major languages have enjoyed some attention, too. Ancient languages, however, are neglected, presumably due to their often high complexity and the extensive study and analysis to which they have been submitted by skilled scholars. Yet most texts have not been integrated in annotated corpora; and though databases such as the Perseus project contain large swathes of morphologically and sometimes syntactically annotated text, the process has been driven largely by manual labour; to give an exhaustive list is not appropriate here, but another such example which is relevant is the PROIEL project, which is also a treebank, i.e. a database of syntactically annotated sentences. It contains data for Herodotus and the New Testament.

On the other hand, there are also corpora which have been tagged using NLP techniques, whose relevancy for this thesis is high and that I have thus described in the next section on methodology.

1.3 METHODOLOGY

In this thesis, we have been largely inspired by two articles by H. Dik and R. Whaling, [Dik and Whaling, 2008 and 2009], in which they document their method for semi-automatically tagging the Perseus project's texts under their own framework, PhiloLogic. They start with a database of analysed forms and a series of tagged texts which they use as initial data to train a decision tree tagger, TreeTagger, developed by Helmut Schmid at the University of Stuttgart, a tool which despite being developed in 1995 has aged well as far as performance is concerned. They achieved remarkable accuracy: with refinements to the training data they achieved 96.2% accuracy during tests on the original training data and 91% accuracy on new data, a result which compares quite favorably when compared to TreeTagger's 97% accuracy when used on German newspaper articles considering the high complexity of ancient Greek and the variety of styles of ancient Greek literature.

It occurred to me that this might be a great method for processing the papyri.info database with a relatively small effort for a high payoff; using data from the Perseus and PROIEL projects, it could be possible to train TreeTagger for both morphology and syntax, apply the resulting parameters to the corpus and thus for the greatest part obviate the need for manual tagging. Given the extent of the corpus (about 50,000 texts containing almost 4,500,000 words), achieving even 85% accuracy would reduce the amount of untagged words to 675,000, many of which I would expect to be proper names or morphologically 'erroneous' forms as are often found in the papyri, data which could itself be analysed with regular expressions and then used to improve the training data.²

² As I set out to verify the originality of my thesis, I found that this statistical approach has been used before for textual criticism! *Vide* Mimno and Wallach [2009], an abstract of which may be found at <http://people.cs.umass.edu/~wallach/publications/mimno09computational.txt>.

2 | THE CORPUS

CONTENTS

2.1	Goals	9
2.2	Design	10
2.2.1	Initial inquiries	10
2.3	Technical	12
2.3.1	Requirements	12
2.3.2	Extracting training data	12
2.3.3	Principles of statistical natural language processing	12
2.3.4	Accuracy testing	12

2.1 GOALS

Our objectives for the modified corpus is threefold. Firstly, we want to morphologically annotate it with reasonable accuracy; secondly, we also want to add syntactical annotation; thirdly, we want to ensure the corpus is compatible with a broad range of tools.

The former two goals are the central enterprise of this thesis. Morphological annotation will pose little theoretical problems, as the morphological features of Greek are not subject to discussion; but for syntactical annotation, a variety of theoretical frameworks are in existence. The fact of the matter is that I have chosen to take the route which is technically easiest: to use the tenets of dependency grammar. The reason Leuven : K.U.Leuven. Faculteit Letteren, 2011 is twofold: on one hand, our training data, the New Testament annotated by PROIEL, itself is structured using a dependency model; on the other hand, dependency trees, in contrast to constituent analysis, assigns each word or morph one node, which is easier to ‘digest’ for a computer program.

Our third goal is to format the corpus in a way that will ensure broad compatibility. This is desirable for several reasons. A first is that a broadly used format offers many options for data treatment; for instance, comma-separated text files can be read by many programs, from the simplest text editor to the most complex software packages, while XML can very easily be integrated in a web interface or parsed and transformed by a variety of programs. A second is that corpus is to be published on GitHub for collaborative editing, or possibly pulled into the main papyri.info repository. This depends on our first point: the idea of GitHub is that all projects are open source and can be contributed to by others, which is facilitated by using a format with wide adoption.

2.2 DESIGN

2.2.1 Initial inquiries

Before arriving at the approach described above, I attempted several other possible methods for automatic analysis. A first idea was to integrate the XML files in a PhiloLogic setup. PhiloLogic is a tool developed by the ARTFL project and the Digital Library Development center at the University of Chicago and released under the GNU General Public License. Its original purpose was to serve as the full-text search, retrieval and analysis tool for large databases of French literature.

PhiloLogic has support for TEI XML and boasts an impressive array of features: it can search through text and deliver KWIC results filtered by frequency and metadata, as well as collocation tables and word order information. Furthermore, the development web page cites support for bibliographic backends in MySQL databases, out of the box operation, interoperability across Unix-based systems, etc.

The tool seemed promising; especially the built-in support for XML and Unicode drew my attention, along with the built-in features for linguistic analysis. It essentially would have made my work easier and allowed me to invest more time in investigating a few linguistic topics. In the end, however, the software did not fit my needs in a few respects.

Firstly, PhiloLogic requires Apache, Perl and several CPAN modules, as well as gawk, gdbm, and agrep. Not a huge amount of dependencies, but still less than using Python and the NLTK, which does not require a server or SQL database to run; also a big problem is the out-of-the-box incompatibility of the required CPAN modules with 64-bit systems. Setting PhiloLogic up was not as easy as advertised; I had to resort to a virtual machine running 32-bit Debian Linux until I was able to discover a method that enabled compatibility with Mac OS X 10.7, my OS of choice.

Secondly, XML support includes standard TEI but has issues with EpiDoc; notably, in my experience, headers were treated as text and EpiDoc's complicated tagset was incorrectly rendered; for instance, original readings and corrections were concatenated in the text browser, an undesirable situation if we wish to have any hand in our choice of corpus, and a possible source of statistical errors. Developing brand new XSLT stylesheets to convert EpiDoc to a simpler form of TEI markup or to raw text could perhaps have been a feasible solution; but PhiloLogic seemed to behave erratically in general and a first attempt at feeding the system raw text (advertised as one of its capabilities) turned out a spathe of errors.

Thirdly, PhiloLogic's feature set is extensive, but I soon came to realize that it does not offer much more than papyri.info already does; it is essentially a robust concordancing application through a web interface. I wished not only to have a lemmatised corpus searchable by word proximity, but also by word relations, which would shed more light

on the syntax of the Greek of the papyri, as this is arguably the largest gap in scholarship due to the age of Mayser's *Grammatik* and the incomplete state of Gignac's *Grammar*. Furthermore, the tool as available to the public does not integrate the morphological database of Greek developed for Perseus under PhiloLogic — a regrettable point, really, since it was exactly this capacity I wished to exploit in the first place. Unicode search for Greek also seems to be in need of improvement; the search form still requires transcription. Modifying the tool would have been unfeasible given the size of the source code.

Thus I discarded PhiloLogic from my options and set out to find or develop a simpler system.¹ After the technical struggle with PhiloLogic, I decided to focus my efforts on concordancers that could replicate its functions. A quick consultation of Stanford's list of resources [Stanford NLP Group, 2011] for statistical natural language processing and corpus-based computational linguistics directed me to WordSmith Tools, which is regrettably Windows-only and a commercial, closed source program to boot. I set out on a search for open-source alternatives and found a diversity of programs, most of them without the same functionality as WordSmith, and invariably clunky or antiquated — making WordSmith the only option, but even that did not seem viable for my ends.

I continued my search and stumbled across a very interesting piece of software that is enjoying a good amount of popularity as a didactic tool for computational corpus linguistics; the Python NLTK, short for Natural Language Tool Kit, is not a stand-alone application, but rather a set of “open source Python modules for research and development in natural language processing and text analytics bundled with data and documentation” [Natural Language Toolkit, 2012]. These modules implement diverse functionalities useful for natural language processing and allow their easy integration into Python programs. The Python NLTK seemed to offer a more than solid amount of features: corpus readers, tokenizers, stemmers, taggers, chunkers, parsers, classifiers, clusterers, tools for semantic interpretation and metrics were all integrated from the get-go and are easily combined and extended. The NLTK was developed for English, but its extensibility meant it could be applied to ancient Greek with relative ease. A further advantage is the fact that Python is very well-suited to text manipulation and parsing XML, which is an advantage when working with the papyri.info corpus; it is also cross-platform: Python interpreters are available for all major platforms.

Now that I had a solid tool for interpreting corpora, the only issue remaining was to construct a linguistically annotated corpus out of the papyri.info corpus. The extent of the corpus makes extensive manual tagging by a single person unfeasible; thus I set out to look for an automated tagger which could be used, immediately or with some effort, on ancient Greek. A variety of taggers seem to be available,

¹ I also considered setting up a mirror of papyri.info on a personal server and modifying the search functionality; this turns out not to be a task for the weak-hearted. The required setup causes substantial overhead for our purposes, too much to be a reasonable solution.

but I settled on TreeTagger, developed by Helmut Schmid, for several reasons:

2.3 TECHNICAL

2.3.1 Requirements

To create the annotated corpus as I have, it is necessary to have the following installed (older versions may work but have not been tested):

- a UNIX-like operating system, i.e. Mac OS X, any variety of Linux or BSD (using Windows should also be possible, since all tools used are portable and cross-platform, but the file location hierarchies in some scripts will not work);
- Python 2.7.3, found at <http://www.python.org/> (do not use Python 3.x, the NLTK is not compatible with it);
- Perl 5.x <http://www.perl.org/>;
- Saxon-HE 9.x, found at <http://saxon.sourceforge.net/>;
- Helmut Schmid's TreeTagger, found at <http://www.ims.uni-stuttgart.de/projekte/corplex/TreeTagger/>;
- the Python Natural Language Toolkit (NLTK), found at <http://nltk.org/>;
- SQLite 3, found at <http://www.sqlite.org/>.

All the above, excluding Mac OS X, are freely available.

2.3.2 Extracting training data

2.3.3 Principles of statistical natural language processing

This subsection is intended as an overview of some key concepts on which natural language processing is based; it is based mainly on Manning and Schütze [1999].

Formal languages

Hidden Markov models

Zipf's law

2.3.4 Accuracy testing

3 | THE TOOL

CONTENTS

3.1	Goals	13
3.2	Design	13
3.3	Technical	13
3.4	Interface	14

3.1 GOALS

1. The tool must be able to read EpiDoc XML and a subset of additional tags added for linguistic annotation. It should be, essentially, a text-based replica of the Papyrological Navigator as found at papyri.info.
2. search
3. statistics
4. syntactic trees

3.2 DESIGN

The tool works from the command line; while this type of interface has the obvious disadvantage of being a type interface most people cannot work with right off the bat, command line programs generally have an edge over programs with a graphical user interface in some aspects. They are smaller, faster, more efficient and flexible, and in Unix environments can easily be integrated in a larger workflow; it is possible, for instance, to take the output of a program and pipe it immediately into a text file or into another program for processing, or to write a script which automates the use of the program.

3.3 TECHNICAL

The tool is written in Python 3.

3.4 INTERFACE

The basic command for launching the script is `python tjufy` from within a terminal or simply `tjufy` from within the Python interpreter (from here on I will simply write `tjufy`). Running it without options like this will display a help dialog; the same goes for running `python tjufy -h` or `python tjufy -help`.

4 | APPLICATIONS

4.1 COLLABORATIVE EDITING

4.2 CORPUS-BASED GRAMMARS

4.3 TEXTUAL CRITICISM

5 | CONCLUSION

BIBLIOGRAPHY

Dik, Helma and Richard Whaling

- 2008 “Bootstrapping Classical Greek Morphology”, in *Digital Humanities*. (Cited on p. 7.)
- 2009 “Implementing Greek Morphology”, in *Digital Humanities*. (Cited on p. 7.)

Evans, Trevor V. and Dirk D. Obbink

- 2010 *The Language of the Papyri*, Oxford University Press. (Cited on pp. v, 1, 3.)

Gignac, Francis Thomas

- 1976 *A Grammar of the Greek Papyri of the Roman and Byzantine Periods. I. Phonology*. Milano: Goliardica. (Cited on p. 1.)
- 1981 *A Grammar of the Greek Papyri of the Roman and Byzantine Periods. II. Morphology*. Milano: Goliardica. (Cited on p. 1.)

Kapsomenos, Stylianos G.

- 1938 *Voruntersuchungen zu einer Grammatik der Papyri der nachchristlichen Zeit*, Münchener Beiträge zur Papyrusforschung und antiken Rechtsgeschichte, München: C. H. Beck. (Cited on p. 2.)
- 1957 “Ἐρευνᾶν εἰς τὴν γλῶσσαν τῶν ἐλληνικῶν παπύρων. Σείρα Πρώτη”, *EEThess*, vii, pp. 225–372. (Cited on p. 2.)

Ljungvik, Herman

- 1932 *Beiträge zur Syntax der spätgriechischen Volkssprache*. Vol. 27, Skrifter utgivna av K. Humanistiska Vetenskaps-Samfundet i Uppsala, 3, Uppsala & Leipzig: Humanistiska Vetenskaps-Samfundet. (Cited on p. 2.)

Mandilaras, Basileios G.

- 1973 *The verb in the Greek non-literary papyri*, Athens: Hellenic Ministry of Culture and Sciences. (Cited on pp. 1, 2.)

Manning, Christopher D. and Hinrich Schütze

- 1999 *Foundations of Statistical Natural Language Processing*, Cambridge, MA and London: MIT Press. (Cited on p. 12.)

McCarthy, Michael and Anne O’Keeffe

- 2010 “What are corpora and how have they evolved?”, in *The Routledge Handbook of Corpus Linguistics*, ed. by Anne O’Keeffe and Michael McCarthy, London and New York: Routledge. (Cited on pp. 3, 4.)

Mimno, David and Hanna Wallach

- 2009 "Computational Papyrology (presentation)", in *Media in Transition 6: Stone and Papyrus, Storage and Transmission*, Cambridge, MA. (Cited on p. 7.)

Natural Language Toolkit

- 2012 *Natural Language Toolkit*, <http://www.nltk.org/>. (Cited on p. 11.)

Palmer, Leonard Robert

- 1934 "Prolegomena to a Grammar of the post-Ptolemaic Papyri", *J. Th. S.*, xxxv, pp. 170–5. (Cited on p. 2.)
- 1945 *A Grammar of the post-Ptolemaic Papyri: Accidence and Word-Formation*, London: Oxford University Press. (Cited on p. 2.)

Porter, S. E. and M. B. O'Donnell

- 2010 "Building and Examining Linguistic Phenomena in a Corpus of Representative Papyri", in *The Language of the Papyri*, Oxford University Press, pp. 287–311. (Cited on p. 3.)

Salonius, A. H.

- 1927 *Zur Sprach der griechischen Papyrusbriefe*, vol. i, Die Quellen, Akademische Buchhandlung. (Cited on p. 2.)

Stanford NLP Group

- 2011 *Statistical natural language processing and corpus-based computational linguistics: An annotated list of resources*, <http://www-nlp.stanford.edu/links/statnlp.html>. (Cited on p. 11.)

Turing, Alan Mathison

- 1950 "Computing Machinery and Intelligence", *Mind*, 59, 236, pp. 433–460. (Cited on p. 6.)